

Football Broadcast Player Detection and Re-Identification using YOLOv11

Final Project Report

Arman Shaik

1 Objective

The objective of this project is to design a robust, real-time system capable of detecting and consistently tracking football players in broadcast videos. A primary challenge addressed is ensuring that each player maintains a unique and persistent identity across video frames—even when temporarily occluded or out of view due to camera motion or scene transitions.

2 Approach and Methodology

The following steps were undertaken to build the Football Player Detection and Re-Identification System:

1. **Frame Extraction:** A 15-second football broadcast video was used as the input source. A total of 375 frames were extracted using OpenCV for annotation and training.
2. **Dataset Creation using Roboflow:** Each frame was uploaded to Roboflow, where manual annotations were performed. The labeled categories included:
 - Football players
 - Referee
 - Football

Data augmentation techniques such as flipping, brightness/contrast adjustment, and blurring were applied. The final dataset contained:

- 663 images total
- 612 images for training
- 38 images for validation
- 13 images for testing

The dataset is publicly available on Roboflow at: <https://universe.roboflow.com/roboflow-jvuqo/football-players-detection-3zvbc/dataset/1>

3. **Model Selection and Training:** The YOLOv11 model was selected for object detection. The training process was carried out using the Jupyter Notebook file Football Analysis System (1).ipynb, where the model was trained for 150 epochs with early stopping applied. The best model checkpoint was saved as best.pt.
4. **Post-Training Inference:** Two inference pipelines were implemented:
 - yolo_inference.py – for basic object detection using pretrained weights
 - main.py – for detection and player tracking using the custom-trained model
5. **Re-Identification and Output:** The output included annotated videos with consistent player IDs, and tracking metadata was saved as a .pkl file inside the tracker stubs/ folder.

3 Techniques Tried and Outcomes

- **Pretrained YOLOv11 Inference:** Initially, the yolo_inference.py script was used to perform basic object detection using pretrained YOLOv11 weights. The model was able to detect generic classes (players, ball) but without persistent IDs or tracking across frames.
- **Custom YOLOv11 Training and Re-ID:** A custom dataset was prepared and annotated using Roboflow. The YOLOv11 model was trained for 150 epochs. The training achieved a final mAP@50 of **0.888** and mAP@50-95 of **0.429**, with high precision and recall across classes like football players, referee, and football.
- **Custom Tracking with Re-Identification:** The main.py script implemented a full pipeline using a custom tracker. The model was initialized with the fine-tuned weights, and the player tracking results were written to tracker_stubs/player_detection.pkl. Visual annotations were drawn across all video frames, and a final output video was saved in output_videos/.

4 Challenges Encountered

- **Re-ID Gaps After Frame Exit:** The model struggled to reassign the same identity to players who left the field of view and returned after a few seconds, leading to fragmented tracking.
- **Visual Similarity Confusion:** Players with nearly identical uniforms, especially teammates, caused the model to incorrectly swap or duplicate identities.
- **Blurred or Occluded Frames:** Fast movements or partial occlusions led to degraded visual features, making identity re-matching highly unreliable.

5 Model Performance and Results

The YOLOv11 model was trained for a maximum of 150 epochs with early stopping enabled (patience = 100). Training was halted at epoch 148, with best performance at epoch 48.

Final Evaluation Metrics (Validation Set)

Category	Precision	Recall	mAP@50	mAP@50-95
Football	0.811	0.778	0.747	0.373
Football Players	0.984	0.922	0.953	0.469
Referee	0.918	0.929	0.963	0.445
Overall	0.904	0.876	0.888	0.429

Table 1: Per-Class Detection and Re-Identification Performance

Speed Performance

- Preprocessing Time: 0.2 ms/image
- Inference Time: 17.4 ms/image
- Postprocessing Time: 1.1 ms/image

```

all      11      152      0.87      0.883      0.868      0.373
Epoch  GPU_mem  box_loss  cls_loss  dfl_loss  Instances  Size
147/150  10.4G     0.8305   0.3913   1.048     156        640: 100% 14/14 [00:10<00:00, 1.39it/s]
Class   Images  Instances  Box(P   R      mAP50   mAP50-95): 100% 1/1 [00:00<00:00, 4.31it/s]
all      11      152      0.872     0.883   0.864   0.371

Epoch  GPU_mem  box_loss  cls_loss  dfl_loss  Instances  Size
148/150  10.4G     0.8166   0.3901   1.045     170        640: 100% 14/14 [00:10<00:00, 1.38it/s]
Class   Images  Instances  Box(P   R      mAP50   mAP50-95): 100% 1/1 [00:00<00:00, 4.69it/s]
all      11      152      0.876     0.878   0.86   0.364

EarlyStopping: Training stopped early as no improvement observed in last 100 epochs. Best results observed at epoch 48, best
model saved as best.pt.
To update EarlyStopping(patience=100) pass a new patience value, i.e. `patience=300` or use `patience=0` to disable EarlyStop
ping.

148 epochs completed in 0.492 hours.
Optimizer stripped from runs/detect/train2/weights/last.pt, 51.2MB
Optimizer stripped from runs/detect/train2/weights/best.pt, 51.2MB

Validating runs/detect/train2/weights/best.pt...
Ultralytics 8.3.160 Python-3.11.13 torch-2.6.0+cu124 CUDA:0 (Tesla T4, 15095MiB)
YOLO11l summary (fused): 190 layers, 25,282,396 parameters, 0 gradients, 86.6 GFLOPs
Class   Images  Instances  Box(P   R      mAP50   mAP50-95): 100% 1/1 [00:00<00:00, 4.16it/s]
all      11      152      0.904     0.876   0.888   0.429
football 9        9        0.811     0.778   0.747   0.373
football players 11     131     0.984     0.922   0.953   0.469
refree   7        12      0.918     0.929   0.963   0.445

Speed: 0.2ms preprocess, 17.4ms inference, 0.0ms loss, 1.1ms postprocess per image
Results saved to runs/detect/train2
🔥 Learn more at https://docs.ultralytics.com/modes/train

```

Figure 1: YOLOv11 Training Summary: Loss curves, class-wise metrics, and evaluation output

6 Future Work and Possible Extensions

To enhance performance and robustness, especially in complex real-world settings, the following improvements can be considered:

- **Strengthen Re-Identification Across Gaps:** Train the system to better associate players who momentarily exit and re-enter the frame using more refined embeddings or temporal matching techniques.
- **Develop Player-Specific Re-ID Features:** Tailoring Re-ID by using features like jersey text, body build, or visual patches like logos could make the system more resilient to appearance overlap.
- **Apply Motion-Based Forecasting:** Motion prediction strategies such as Kalman filters or trajectory estimation can help maintain consistent identities during short disappearances.
- **Boost Efficiency for Real-Time Applications:** Convert and optimize the model using inference engines such as ONNX or TensorRT to reduce frame processing delays.
- **Scale to Multi-View Tracking:** Extend the architecture to handle multi-camera input and match players across views using cross-camera Re-ID techniques.

7 Conclusion

This project successfully demonstrated a working pipeline for football player detection and re-identification in broadcast videos using YOLOv11. The system showed significantly better performance when trained on a custom dataset and integrated with a tracking mechanism. Future improvements can include real-time deployment, jersey number OCR, and larger-scale datasets for robust generalization.